



GPU Monitoring with DCGM

Unit 5



Outline

- DCGM Overview
- DCGM CLI Usage and Examples



Objectives

By the end of this unit, you will be able to:

- Define the purpose and features of DCGM
 - Apply the DCGM command-line interface (CLI) to perform common tasks
-

DCGM Overview



NVIDIA Data Center GPU Management (DCGM)

- The NVIDIA® Data Center GPU Manager (DCGM) is an intelligent, lightweight user space agent that provides monitoring capabilities of NVIDIA GPUs in the cluster.
- DCGM simplifies administration by performing a variety of functions related to the GPUs available on the system, such as:
 - GPU behavior monitoring
 - GPU configuration management
 - GPU policy oversight
 - GPU health and diagnostics
 - GPU accounting and process statistics
 - NVSwitch configuration

DCGM on a DGX System

Focus Areas:

- Manage GPUs as collections of related resources
- Configure NVSwitch
- Define and enforce GPU configuration state
- Automate GPU management policies
- Provide robust, online health and diagnostics
- Enable job-level statistics and accounting

DCGM Framework Includes:

- Programmatic access via C and Python
- Python interfaces for admin-centric scripting environments
- CLI-based tools to provide an interactive out-of-the-box

DCGM CLI Usage and Examples



DCGM CLI

- Standalone Mode is the NVIDIA preferred mode of operation
- DCGMI is the CLI interface for the standalone mode of DCGM
- Allows users to easily configure, manage and monitor GPUs and interconnections
- For more details on the DCGM usage refer to the user guide:
- <https://docs.nvidia.com/datacenter/dcgm/latest/dcgm-user-guide/index.html>

DCGM Help Page

- Use '**dcgmi**' to run DCGM CLI commands
- Run '**dcgmi --help**' command for details
- Subsystems are primary method used to interact with DCGM CLI

```
~ $ dcgmi --help
Usage: dcgmi
      dcgmi <subsystem>
      dcgmi -v

Flags:
  -v   vv          (OR required)  Get DCGMI version information
      subsystem    (OR required)  The desired subsystem to be accessed.
                           Subsystems Available:
...
Please email dcgm-support@nvidia.com with any questions, bug reports, etc.

NVIDIA Datacenter GPU Management Interface

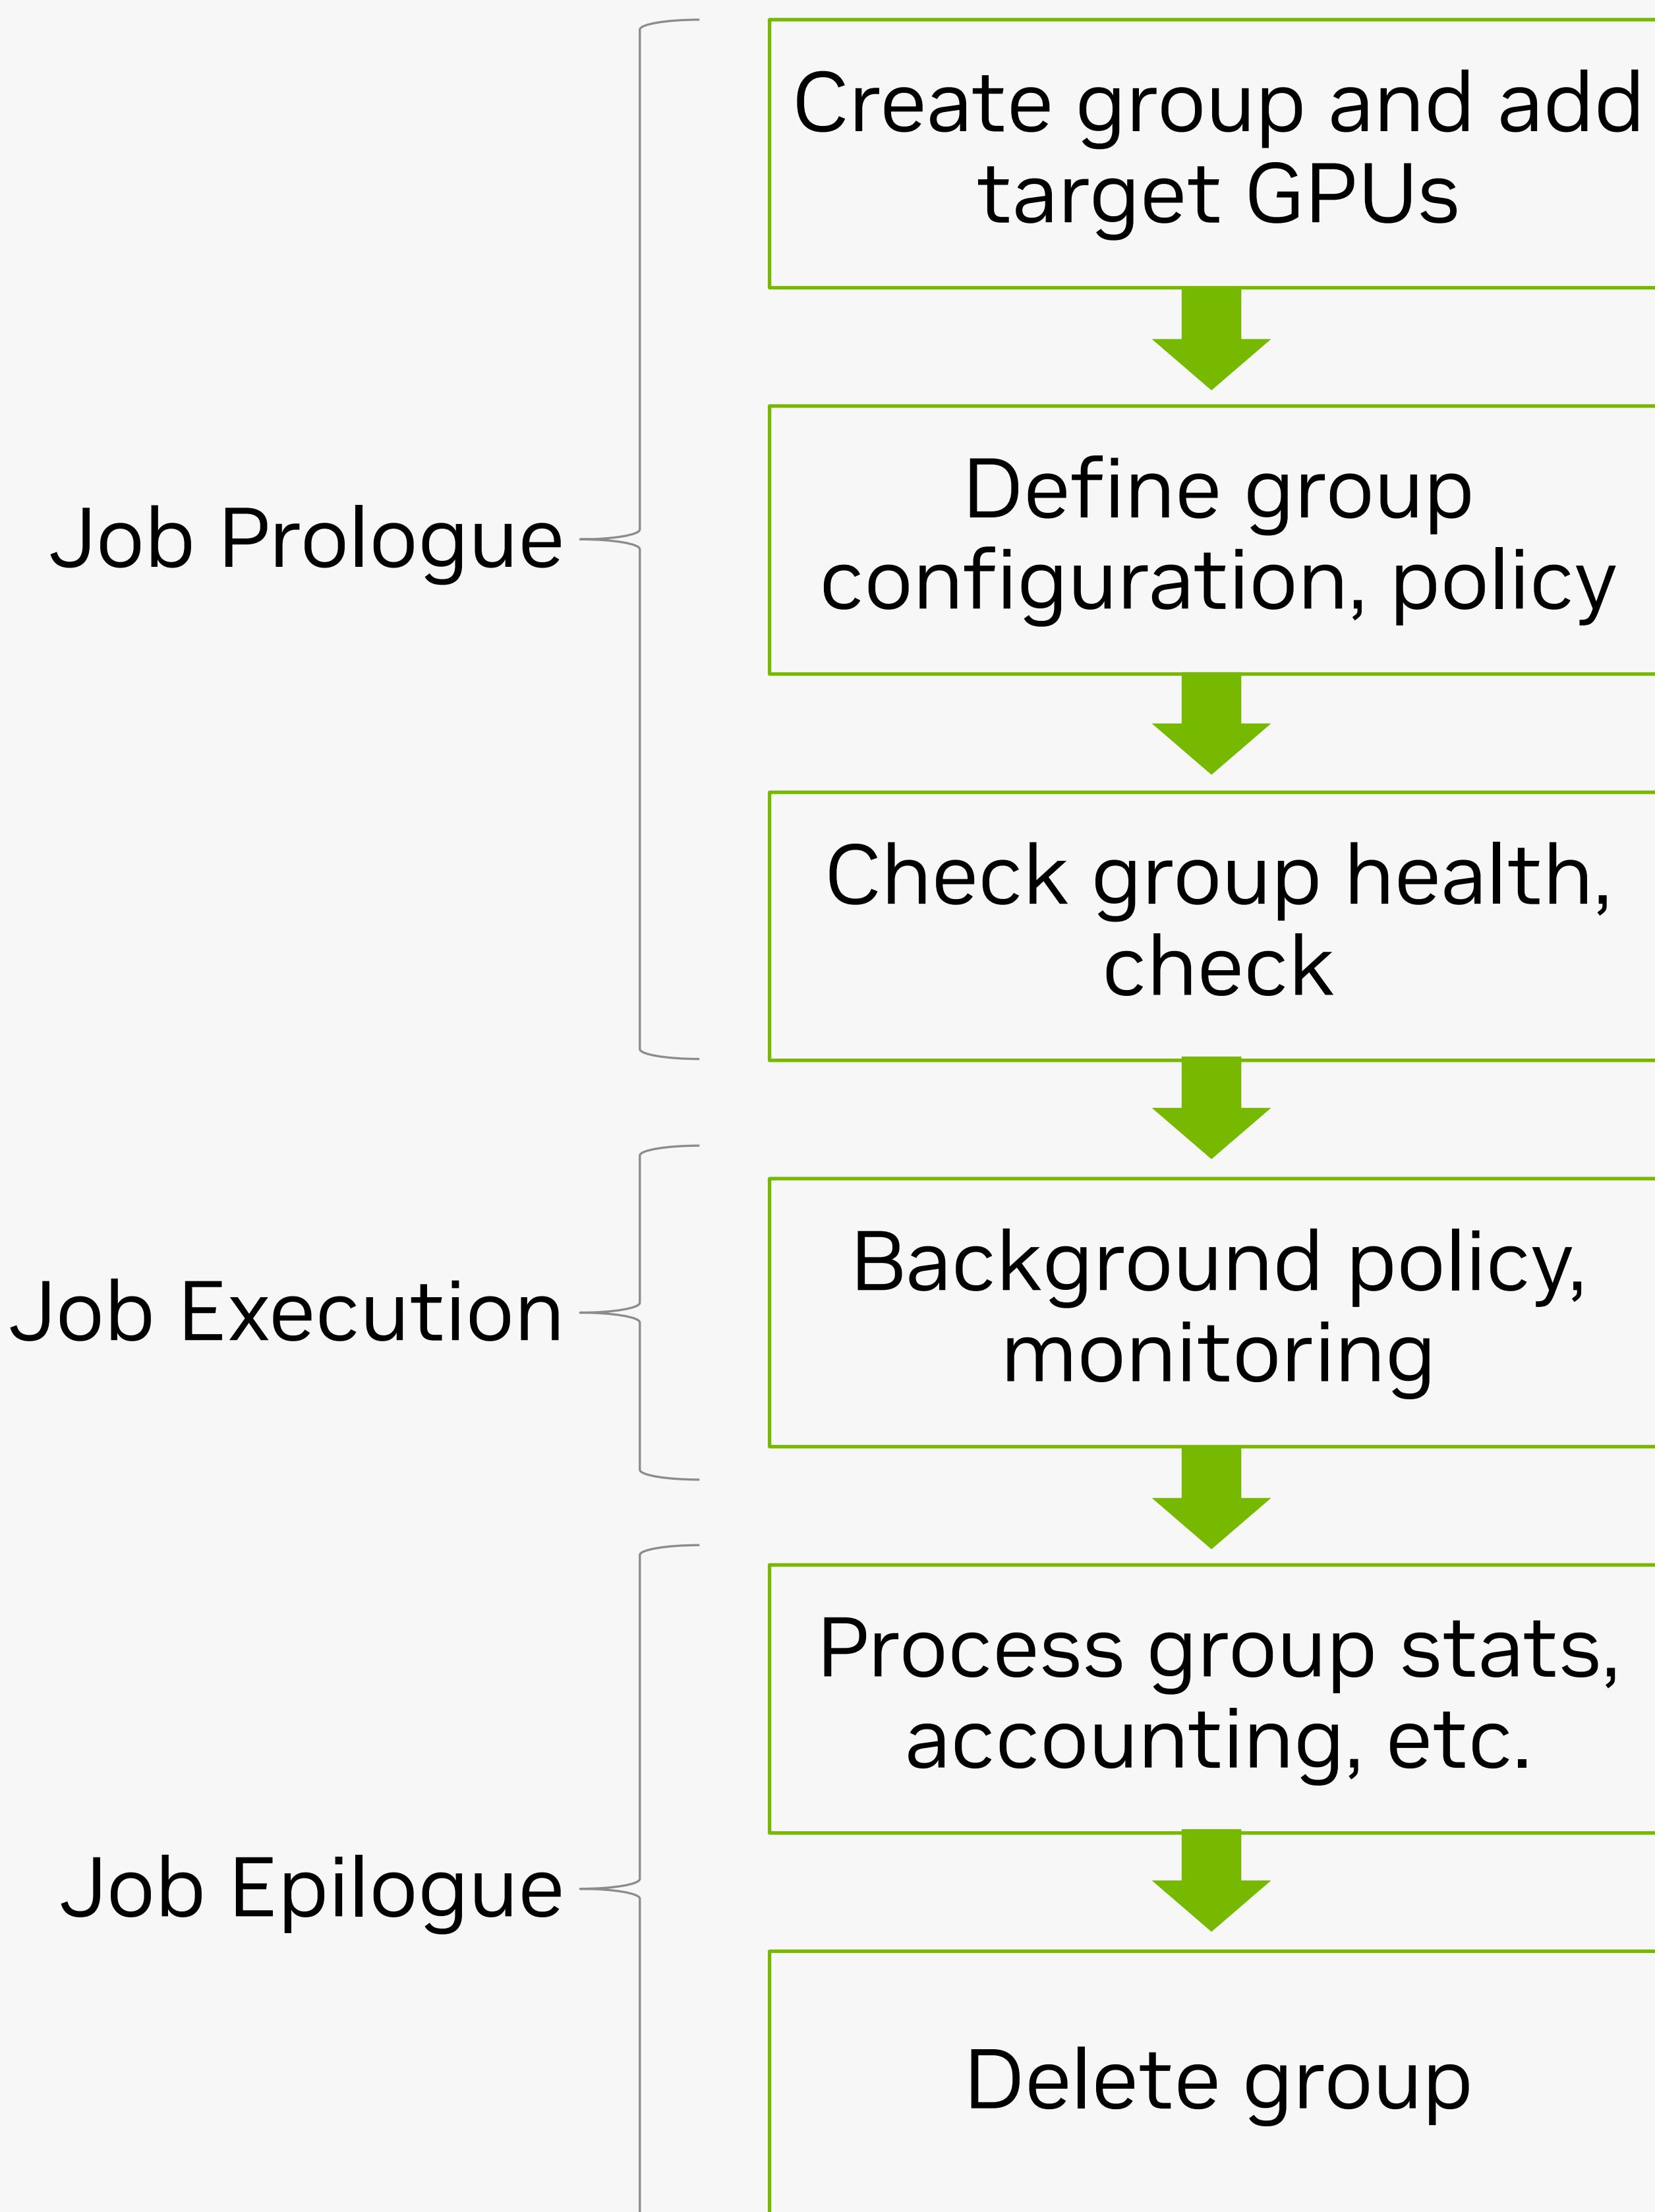
...[output truncated]
```

DCGM CLI Subsystems

Subsystem	Description
set	Configure hostengine settings
profile	Control and list DCGM profiling metrics
modules	Control and list DCGM modules
dmon	Stats Monitoring of GPUs [dcgmi dmon -h for more info]
nvlink	Displays NvLink link statuses and error counts [dcgmi nvlink -h for more info]
introspect	Gather info about DCGM itself [dcgmi introspect -h for more info]
discovery	Discover GPUs on the system [dcgmi discovery -h for more info]
group	GPU Group Management
config	Configuration Management
health	Health Monitoring
policy	Policy Management
diag	System Validation/Diagnostic
stats	Process Statistics
topo	GPU Topology

DCGMI CLI Usage

- Almost all DCGM operations take place on groups
- Groups are intended to help the user to manage collections of GPUs as a single abstract resource
- Setting up a group of GPUs for a job would follow the flowchart



DCGMI Creating and Deleting a Group

- A group can be created using ‘`dcgmi group -c <Group_Name>`’
- To get a listing of existing groups, use ‘`dcgmi group -l`’
- To delete an existing group, use ‘`dcgmi group -d <Group ID>`’

```
~ $ dcgmi group -c GPU_Group  
GPU_Group Successfully created group "GPU_Group" with a group ID of 1
```

```
~ $ dcgmi group -l  
1 group found.  
+-----+  
| GROUPS  
+=====+  
| Group ID | 1  
| Group Name | GPU_Group  
| GPU ID(s) | None  
+-----+
```

```
$ dcgmi group -d 1  
Successfully removed group 1
```

DCGMI Adding GPUs to a Group

- DCGMI discovery can be used to find a list of available GPUs
 - Selected GPUs can be added to an existing group

```
~ $ dcgmi discovery -l
2 GPUs found.

+-----+
| GPU ID | Device Information
+=====+
| 0      | Name: Tesla K80
|         | PCI Bus ID: 0000:07:00.0
|         | Device UUID: GPU-00000000000000000000000000000000
+-----+
| 1      | Name: Tesla K80
|         | PCI Bus ID: 0000:08:00.0
|         | Device UUID: GPU-11111111111111111111111111111111
+-----+
```

DCGMI GPU Configuration & Policy

- Use the **config** subcommand to set the target configuration for the group
- The **policy** subcommand is used to set up automatic GPU behaviors (e.g., action, notification) based on various conditions.
- The **--set** and **--get** flags are used to set and read configuration or policy information
- The **--help** flag gives the details on the settings

```
~ $ dcgmi config -g 1 --set -c 2
Configuration successfully set.
```

```
~ $ dcgmi policy -g 2 --set 0,0 -p
Policy successfully set.
```

DCGMI Enable and Check Stats

- DCGMI stats must be enabled to track GPUs status
- Status can be tracked across jobs and groups of GPUs

```
~ $ dcgmi stats -g 2 --enable
Successfully started process watches.
```

```
~ $ dcgmi stats --pid 1234 -v
+-----+
| GPU ID: 0 |
+=====+=====+
| ----- Execution Stats ----- |
| Start Time * | Tue Nov 3 17:36:43 2015
| End Time * | Tue Nov 3 17:38:33 2015
| Total Execution Time (sec) * | 110.33
| No. of Conflicting Processes * | 0
+----- Performance Stats -----+
```

[output truncated]

(*) Represents a process statistic. Otherwise device statistic during process lifetime listed.

Summary

Now that you have completed this unit, you should be able to:

- Define the purpose and features of DCGM
- Apply the DCGM command-line interface (CLI) to perform common tasks

